



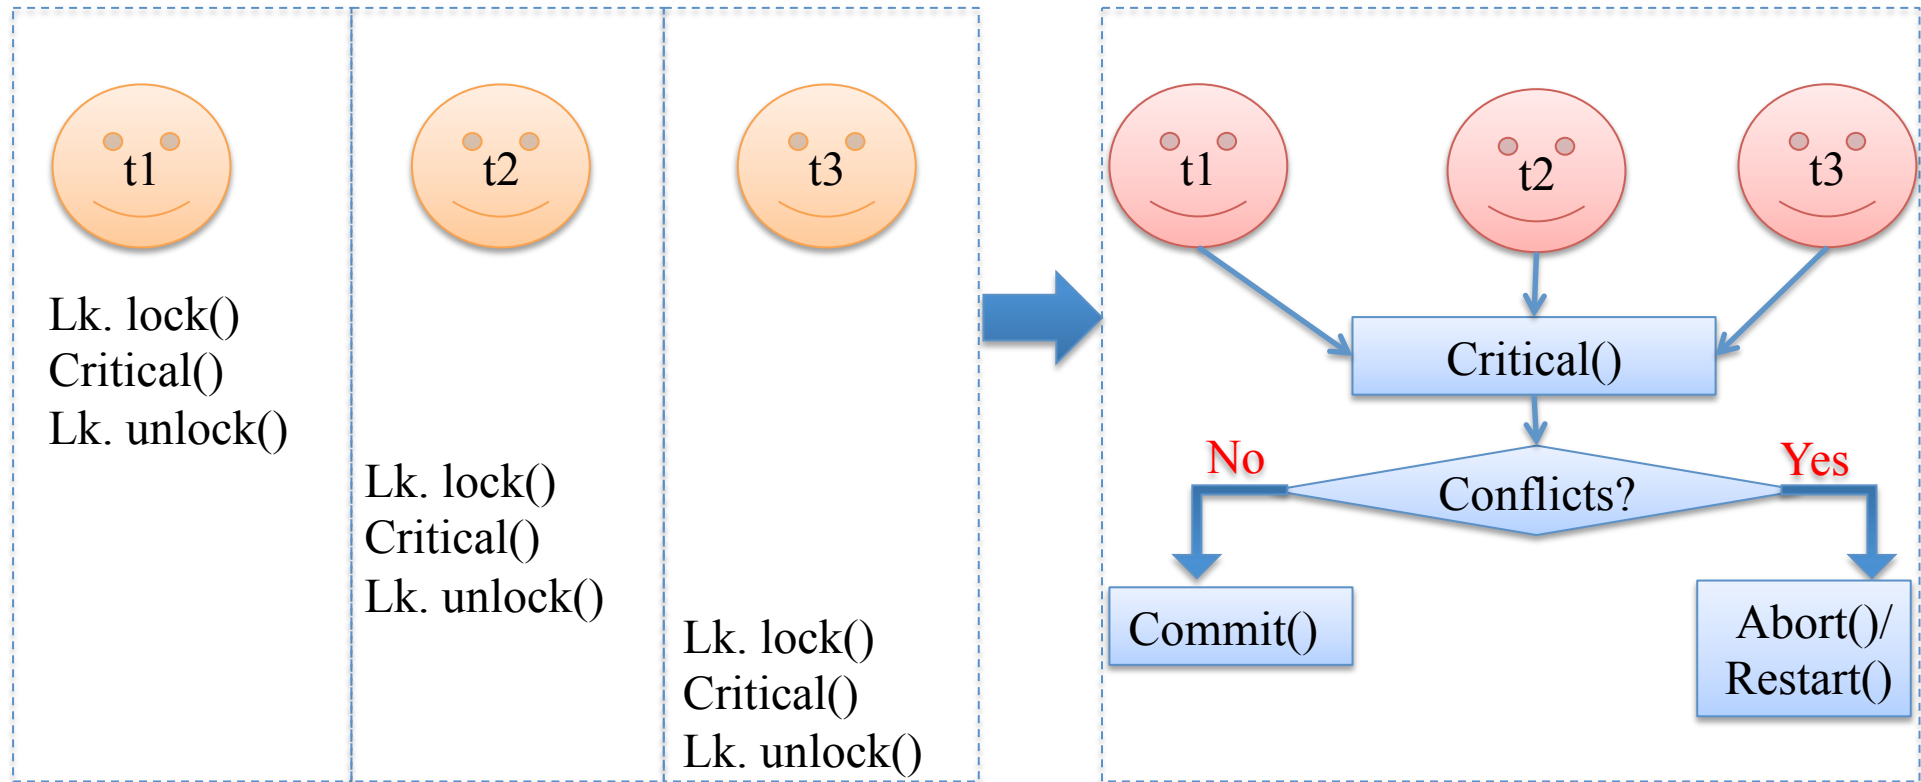
# On the Platform Specificity of STM Instrumentation Mechanisms

**Wenjia Ruan**, Yujie Liu, Chao Wang, and Michael Spear  
Computer Science and Engineering Department  
Lehigh University



# Background: Transactional Memory (TM)

- Mechanism to simplify concurrent programming
- Replace critical sections with transactions
- Transactions use speculation under the hood






# Implementing Software TM (STM)

Programmer writes code:

```
__transaction_atomic {  
  if (B == 0)  
    A ← 1  
  else  
    B ← 0  
}
```



Compiler generates code:

```
TM_CHECKPOINT()  
TM_BEGIN()  
if (TM_READ(B) == 0)  
  TM_WRITE(&A, 1)  
else  
  TM_WRITE(&B, 0)  
TM_COMMIT() // may fail  
             // and cause  
             // txn to  
             // retry
```

- Many possible implementations of the **TM\_** functions



# TM Availability and Use

- Latest product announcements from Intel and IBM include **Hardware TM** (HTM) support
  - Not all transactions can be executed using HTM (e.g., big transactions, long-running transactions)
  - HTM can enter pathological situations (e.g., livelock)
  - Software fallback is required
  - Lots of legacy systems would still need STM if they want to use transactional memory
- **Software TM** progress
  - C++ standardization effort underway since 2009
  - Support in Scala, Clojure, Haskell



# This Paper

- How does the **platform** (OS, CPU) affect the manner in which a compiler should instrument code to achieve efficient STM?
  - Focus on platform characteristics and relationship to system Application Binary Interface (ABI)
  - Goal: to see what implementation details have platform-specific overheads
- Evaluations on 4 platforms:
  - ARM/Linux
  - SPARC/Solaris
  - IA32/Linux
  - IA32/MacOS



# Main Results

- Many sources of overhead are platform-specific:
  - Accessing thread-local storage (TLS)
  - Mechanisms for adaptivity/autotuning
  - Memory consistency model / memory fences
- Achieving optimality on every platform is a code maintenance nightmare
  - Macros to generate function signatures
  - Extend compiler toolchain to produce function bodies for adaptive libraries
  - Best algorithm for one platform can be worst on another platform
    - We designed a new STM algorithm “Cohorts”



# Generality

- These issues are not specific to STM
  - NUMA-aware algorithms often require TLS access
  - Autotuning and adaptive libraries are common in HPC
  - The world isn't all x86/Linux



# Evaluating TLS Overheads

- STM implementations use **per-thread metadata** (e.g., “transaction descriptors”) to store read/write sets, rollback values, etc.
  - Accessed on every shared memory load/store
- Common ways to access per-thread metadata:
  - Access thread-local storage (**TLS**). (E.g., use `__thread` modifier) --- used currently by GCC compiler.
  - Compiler explicitly manage descriptors. (e.g., the descriptor is passed to the TM library as an **additional parameter**) --- used currently by Oracle TM compiler.
    - Requires changes to function signatures






# Evaluating TLS Overheads

- We modified RSTM to support both conventions
- Executed stress-tests, where there is no work between transactions
  - Red/black tree, equal mix of insert and remove operations, 8-bit values, 50% update transactions
  - Use gcc in all cases



# Evaluating TLS Overheads




(a) Platform: IA32(12-way)/Linux, Algorithm: TinySTM

IA32/  
Linux



**No noticeable difference**

```
get_tls:
mov  %gs:0xffffffff0, %eax
```




(b) Platform: IA32(4-way)/MacOS, Algorithm: TinySTM

IA32/  
MacOS



**5% slowdown**

No OS-level support for \_\_thread...  
Must use pthread\_getspecific()




(c) Platform: SPARC(64-way)/Solaris, Algorithm: TinySTM

SPARC/  
Solaris



**3% slowdown**

```
get_tls:
sethi %hi(0), %g1
xor   %g1, -4, %g1
ld    [%g7 + %g1], %o0
```



(d) Platform: ARM(4-way)/Linux, Algorithm: NOrec

ARM/  
Linux



**15% slowdown!**

```
get_tls:
ldr   r3, [pc, #12]
push  {lr}
bl    __aeabi_read_tp
...
```

**7 instructions,  
Including a  
branch**

Time is  
Energy!



# Summary of TLS costs

- **Performance:**
  - On some platforms, `__thread` has noticeably more latency, on others it is faster
- **Software engineering / code maintenance:**
  - Optimal solution requires different interfaces to a library depending on the architecture
  - Use `#defines` for every function in a library



# Accessing the TM Library

There is no single TM implementation suitable for all circumstances.

- At least limited **adaptivity** is advisable
  - C++ spec needs it
  - Changing mode/algorithm for performance
- Even more likely when HTM arrives




# Accessing the TM Library

Three common approaches for supporting adaptivity within a TM implementation:


1. **Coarse:** use **global function pointers**, which coordinate all transactions' behavior without incurring additional branching overhead in the common case
2. **Fine:** use **per-thread function pointers**, so each thread can make fine-grained decisions based on previous access patterns
3. **Switch:** use conditionals (typically a “**switch**” statement) to globally coordinate which function to call




# Accessing the TM Library




(a) Platform: IA32(12-way)/Linux, Algorithm: TinySTM



(b) Platform: IA32(4-way)/MacOS, Algorithm: TinySTM



(c) Platform: SPARC(64-way)/Solaris, Algorithm: TL2



(d) Platform: ARM(4-way)/Linux, Algorithm: NOrec

## Coarse

- Global function pointers
- Avoid TLS overhead, prevents branching in the common case where adaptivity is not occurring.


```
tmread = LSA.read;  
tmwrite = LSA.write;  
tmcommit = LSA.commit;
```

Optimal on IA32/MacOS:


- TLS is expensive.
- Indirection of function pointers does not incur a noticeable cost




# Accessing the TM Library




(a) Platform: IA32(12-way)/Linux, Algorithm: TinySTM



(b) Platform: IA32(4-way)/MacOS, Algorithm: TinySTM



(c) Platform: SPARC(64-way)/Solaris, Algorithm: TL2



(d) Platform: ARM(4-way)/Linux, Algorithm: NOrec

## Fine

- Per-thread function pointers
- Affords maximum flexibility
- Incurs TLS overhead to locate function pointers

```
*(threads[i]->my_tmread) = (void*)LSA.read_ro;  
*(threads[i]->my_tmwrite) = (void*)LSA.write_ro;  
*(threads[i]->my_tmcommit) = (void*)LSA.commit_ro;
```

Best for IA32/Linux and SPARC/Solaris:  
[IA32]


- Indirect calls and TLS are inexpensive.
- Even better than “None” (avoids branches to recover more performance than it costs).

[SPARC]


- Greatly benefits from the reduction in branches




# Accessing the TM Library




(a) Platform: IA32(12-way)/Linux, Algorithm: TinySTM



(b) Platform: IA32(4-way)/MacOS, Algorithm: TinySTM



(c) Platform: SPARC(64-way)/Solaris, Algorithm: TL2



(d) Platform: ARM(4-way)/Linux, Algorithm: NOrec

## Switch

- Use a “switch” statement to choose the right instrumentation.
- No TLS or indirect call overhead
- Should be fast on architectures with inexpensive indirect branches

```
void tmwrite(addr, val) {  
    switch (alg) {  
        case LSA: LSAWrite(addr, val); break;  
        ...  
    }  
}
```

Optimal on ARM/Linux:

- TLS is expensive.
- Jump tables have low latency





# Summary of Library Access Costs

- Nuanced interaction between TLS and mode-switching
  - More factors than just the cost of TLS
- Not specific to TM library, likely to affect any adaptive or auto tuned library



# Algorithm Selection and Optimization

- Architecture influences which algorithms are worth considering in an auto tuned STM


E.g.:

- Number of cores
- Cost of CAS instructions
- Cost of Memory fences on each read/write
  - Usually ignored with x86 and SPARC
  - but significant on ARM and POWER



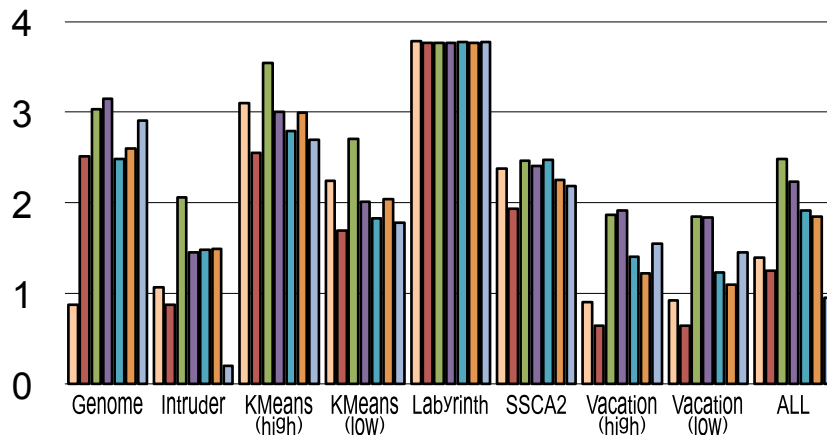
# Evaluation on IA32/Linux and SPARC/Solaris

STAMP on IA32/Linux (12 threads)



1. **LSA**, **TML** and **TLRW** are the only algorithms in which locks are acquired before commit time.
2. **TML**, **OrecELA** and **NOrec** are the only algorithms that comply with the C++ memory model.

STAMP on SPARC/Solaris (64 threads)



- **LSA** offers best performance
- **OrecELA** offers best performance with C++ compatibility



**These are already known.  
What about on ARM?**

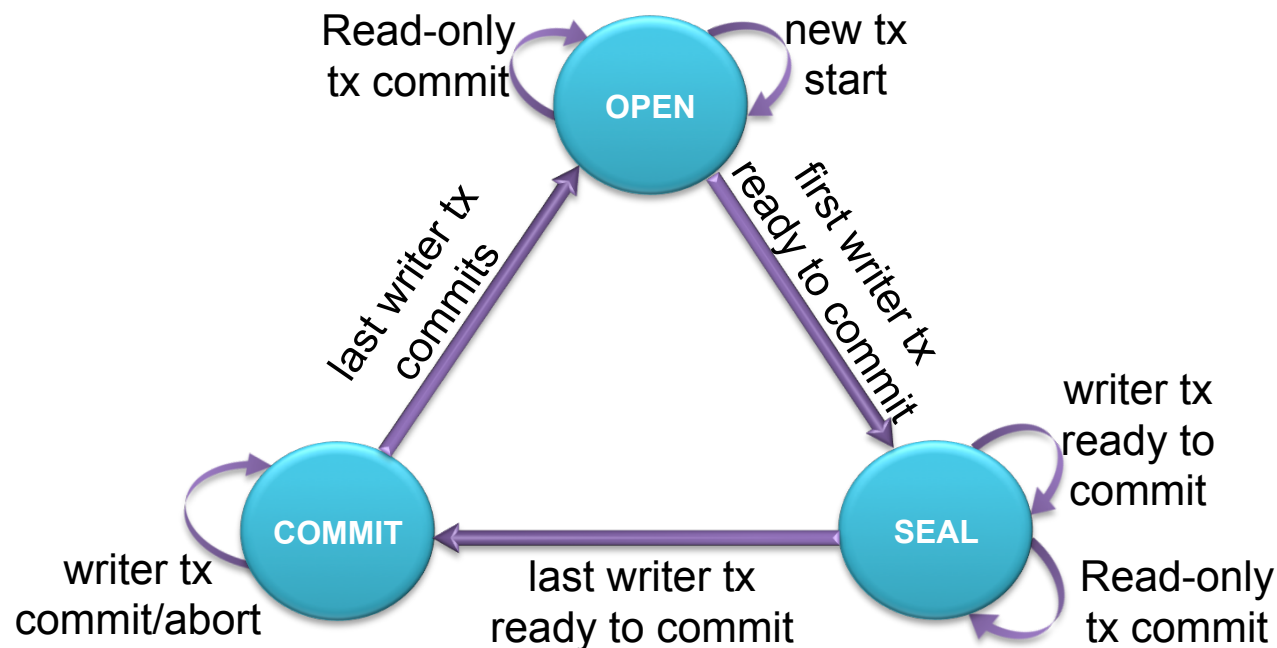


# Algorithm Selection on ARM

- **Does ARM represent a unique design point?**
  - LSA and TL2: each load requires **two** fences.
  - NOrec, TML and OrecELA: **one** fence is required after each load.
  - TLRW requires **one** fence on every load and store.
  - **Fences are expensive!**
- Our new algorithm Cohorts requires **zero** fences load/store
  - Only good for low-core counts
  - Only beneficial due to eliminating fences
  - Performs badly on x86 and SPARC



# Cohorts: A New STM Algorithm for ARM



- Three phases, transactions progress as a cohort
  - Once one writer ready to commit, none can start
  - Once all writers ready to commit, commit one at a time
  - Once all writers finish commit, everyone can start again
- Memory is immutable during execution, no fences needed



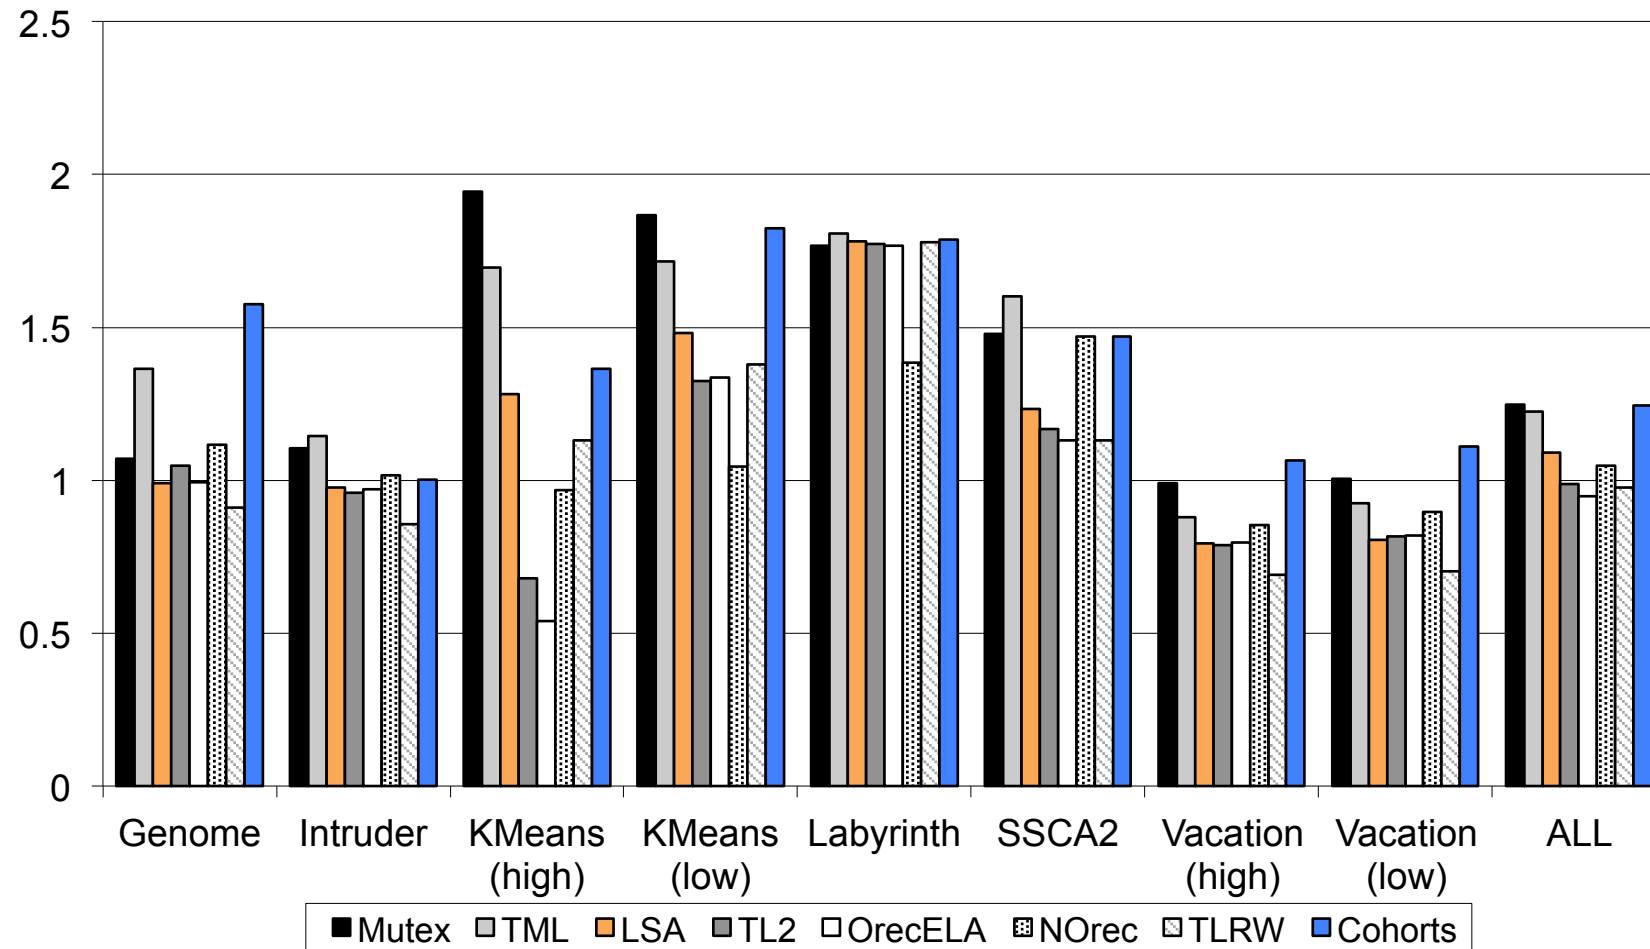
# Cohorts: Simplified TM\_READ()

```
TL2_TM_READ(addr) {  
    // check if exist in write set  
    if (addr  $\in$  writes)  
        return writes[addr];  
  
    // check metadata  
    o1  $\leftarrow$  orecs[addr].getValue();  
    FENCE;  
    // get value of the address  
    v  $\leftarrow$  *addr;  
    FENCE;  
    // check metadata  
    o2  $\leftarrow$  orecs[addr].getValue();  
  
    // validate  
    if (o1 = o2 and o2  $\leq$  start) {  
        reads  $\leftarrow$  reads  $\cup$  {addr};  
        return v; }  
    else ABORT();  
}
```

```
Cohorts_TM_READ (addr) {  
    // check if exist in write set  
    if (addr  $\in$  writes)  
        return writes[addr];  
  
    // get value of the address  
    v  $\leftarrow$  *addr;  
    reads  $\leftarrow$  reads  $\cup$  {<addr, v>};  
    return v;  
}
```



# Evaluation on ARM/Linux (Tegra3: 4 threads)



- Cohorts gives best overall performance
- Still not significantly better than Mutex



## Cohorts: Future Work

- Many cohort-specific compiler optimizations are possible
- Possibly a good fit for Xbox?
- But HTM for ARM would be a better approach





# Conclusions

- Paper focuses on engineering issues
  - Demonstrates that run-time costs of library interface are platform-specific
  - Implications for any multi-platform concurrent or autotuned library
  - Software engineering nightmare to try to provide best alternative in every case
- Some easy fixes (MacOS needs better TLS), but not others
  - Cohorts: a first attempt for STM algorithm on ARM
  - Worst case: need to rethink concurrent algorithms in ARM version of the code



# Thank you ...

- Open sourced, code available at Google Code repository:
  - <http://code.google.com/p/rstm/>
- Contact Info:
  - **Wenjia Ruan:** [rwj@lehigh.edu](mailto:rwj@lehigh.edu)
- Questions?